

# EP2420 Project 2 - Forecasting Service Metrics

André Silva

November 26, 2020

## Task I

In this task we use Linear Regression [2] to predict, not the current state of a service metric, but the future states of said metric.

First, we pre-process the trace, standardizing along each column and removing outliers that deviate more than  $40\sigma$ . We also apply tree-based feature selection, using 10 estimators, to get the top 16 features.

After this, we split the trace into training and test sets, and then transform them into new datasets that follow the structure  $([x^{(t-l)}, \dots, x^{(t)}], [y^{(t)}, \dots, y^{(t+h)}])$ . As we are using analyzing *KV periodic* [1], our step size (i.e. time between each sample of a sequence), is 1 second.

We then perform an experiment to study the impact of  $l$  on the accuracy of a model and its predictions of 0...10 steps into the future by building 11 different Linear Regression [2] models with  $l \in \{0, \dots, 10\}$ . The results can be found below in Figure 1.




Figure 1: Heat-map of NMAE for each Linear Regressor with  $l \in \{0, \dots, 10\}$  when predicting  $y^{(t+h)}$

Figure 1 shows us the NMAE of each model when predicting  $y^{(t+h)}$ . Rows represent the time horizon  $h = 0, \dots, 10$  and columns represent the lag  $l = 0, \dots, 10$ . The color gradient of the heat-map allows us to visualize and extract conclusions better than if we used a normal table.

When we analyze the results by looking at each row we notice that, in most cases, the NMAE decreases as we increase  $l$ . On the other hand, when we look at each column, we see that the NMAE increases as we increase  $h$ .

Both these events are, intuitively, expected. When we increase  $l$ , we are providing more information, or context, to the model, so we expect better predictions. When we increase  $h$ , we are asking for a value that is more distant in time, and so, harder to predict.

This being said, we can also see that, when predicting  $h = 1, \dots, 7$ , the NMAE don't monotonically decrease, when we increase  $l$  from 0, as expected. An explanation to this might be the fact that we built 11 different models, as opposed to  $11^2$  models, and so  $y^{(t+h)}$  predictions are influenced by the noise of unused predictions. Another possible explanation could be that the current state of the system ( $l = 0$ ) can better predict the near future ( $h > 0$ ) than the actual current state ( $h = 0$ ), as system metric variations only have an impact on users after a certain time.

## References

- [1] F. S. Samani, H. Zhang, and R. Stadler. "Efficient Learning on High-dimensional Operational Data". In: *2019 15th International Conference on Network and Service Management (CNSM)*. 2019, pp. 1–9. DOI: 10.23919/CNSM46954.2019.9012741.
- [2] `sklearn.linear_model.LinearRegression`. [https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LinearRegression.html#sklearn.linear\\_model.LinearRegression](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html#sklearn.linear_model.LinearRegression). accessed: 2020-11-26.